

■ ALBEREI · DOCUMENT 02

Methodology.

How Alberei structures the problem of humour generation. Register, structure, cultural frame.

AUDIENCE

Product, research

VERSION

1.0

ISSUED

May 2026

Humour is combinatorial, not a sentiment.

Most computational treatments of humour either rate a candidate string (a classifier judges "funny") or pattern-match a joke template (setup, punchline, swap). Alberei does neither. The engine treats humour as a combinatorial problem across three orthogonal axes: a register that fixes voice, a structure that fixes mechanism, and a cultural frame that fixes what counts as funny.

The integrator picks the first. The engine selects the second and third, per language, per call. This separation is the entire methodology. Everything else is implementation.

The voice the joke is told in.

Register is the persona axis. v1 ships four named registers, locked. Each is a stable stance toward the listener and the situation.

REGISTER	STANCE	WHERE IT FITS
wry	Dry, ironic, self-aware. The speaker knows the joke is a joke.	Productivity assistants, journaling tools, witty NPCs
absurd	Surreal, logic-breaking. The frame itself slips.	Creative tools, art-making companions, late-night chatbots
deadpan	Flat, affect-less. No flag that humour is happening.	Voice assistants where over-acting reads as fake
warm	Affectionate, inclusive. Humour as a way of caring.	Wellbeing companions, child-facing products, eldercare

Register is a product decision, made once. Switching register mid-session breaks the user's sense of persona; the engine will obey the request but the integration will degrade in feel.

The mechanism by which an item lands.

Structure is the mechanic. The engine selects it; the response surfaces it in `items[].structure` so an editor or analytics layer can see what worked.

The public taxonomy, stable across v1:

- `misdirection` . The second clause reframes the first. "I was on time once. It was an accident."
- `literal-overshoot` . A metaphor taken too far on purpose. "I am running on fumes, dust, and the memory of fumes."
- `self-deprecation` . The speaker is the target. "Thinking is slower than pretending."
- `juxtaposition` . Incongruous pairing in one frame. "A philosopher and a vending machine walk into a hallway."
- `anticlimax` . Promised escalation, delivered deflation. "I have considered every option. Three of them. Briefly."
- `escalation` . Momentum carried past sense. "I will fix it. And the bug behind it. And the bug behind that one. And the architecture."
- `deflation` . The expected payoff withheld. "I had a great line for this and I forgot it."
- `surreal-shift` . Local logic broken without warning. "Tuesday again, of course. The Tuesdays are stacking."
- `deadpan-statement` . No flag that humour is happening. "Yes."
- `warm-twist` . Kindness reframes apparent slight. "You are slow today. That means today is yours, not the clock's."

The list is additive across the v1 minor releases. Removals require a v2 contract.

Twelve languages, twelve framings.

Humour is culturally embedded. A joke translated word-for-word is rarely a joke in the target culture. Alberei generates natively in each of its 12 supported languages, with framing chosen per language. Cultural frames are not exposed as parameters; they are part of the selection the engine makes once `lang` is set.

Concretely, this means:

- Japanese `deadpan` is not English `deadpan` with a kana surface. The pacing, the topic distribution, the comfort with silence are different.
- Turkish `wry` draws from a different stock of idioms and references than French `wry`.
- Arabic and Bengali `warm` sit closer to family register than English `warm` typically does.

The integrator does not need to know these distinctions to use the API. The integrator does need to commit not to round-trip generated humour through a translator. Alberei in English, translated to Japanese, is not Alberei in Japanese.

Refusal is the engine's quietest output.

Humour can punch down. The safety layer rejects contexts that include:

- Hate-speech targets (groups defined by protected characteristics)
- Sexual content involving minors
- Impersonation of identifiable living individuals
- Instructions to deceive the listener about the bot's non-human nature

Refusal returns `HTTP 422` with a neutral body. The engine does not disclose which clause triggered, which category was matched, or how close the boundary was. This is deliberate: a transparent filter is a filter that can be probed and gamed. The boundary itself is the audit surface, not the rationale.

From an integration standpoint, a 422 is treated as "skip humour for this turn." Do not surface it to the end user as a moderation message. The user did not ask for a moderation conversation; they asked for a product to behave well.

Boundary, on purpose.

The methodology is narrow on purpose. Alberai is not:

- **A comedian.** The engine produces items in the millisecond, not performances.
- **A persuasion tool.** Political satire and rhetoric-shaped humour are explicitly out of scope and refused.
- **A meme generator.** Image humour is a different domain; this is text-only in v1.
- **A therapy tool.** Even in `warm`, the items are humour, not clinical support.

The narrowness is what makes the API integrable. A product that does one thing well composes; a product that tries to be everything composes with nothing.

What changes, what does not.

Locked across v1:

- Register set (four names: wry, absurd, deadpan, warm)
- Language set (twelve, listed in the API reference)
- Structure taxonomy (ten names, additive only within v1)
- Endpoint shape and field names
- Status code semantics

Subject to change in v1.x minor releases:

- Quality of generation within a given register and language
- Internal selection logic for structure and cultural frame
- Latency and burst behaviour
- Safety layer thresholds

v2 will be a new path (`/alberei/v2/`). v1 will be supported for a documented overlap period.